

Serverless AWS Lambda Explained

what is serverless (lambda in aws) ?

explain in details so that i am able to understand and donot have doubt on this.

Great question, Piyush. **Serverless (AWS Lambda)** is one of those concepts that *sounds confusing* at first, but once it clicks, it actually becomes very logical.

I'll explain it **from zero → internals → real flow → pros/cons → when to use**, so you won't have doubts.

👉 First, clear the biggest myth

✗ "Serverless means there are no servers"

Wrong.

✓ Reality:

- Servers DO exist
- You just don't manage them

AWS manages:

- Servers
- OS
- Scaling
- Patching
- Availability

You manage:

- Only your code

This is why it's called **server-less (from developer's perspective)**.

✳️ What is Serverless?

📌 Definition (simple)

Serverless is a cloud execution model where you run code without provisioning or managing servers, and you pay only when your code runs.

In AWS, the most popular serverless service is:

🔥 AWS Lambda

| Lambda lets you run small pieces of code (functions) in response to events.

🧱 Traditional Server-based App vs Serverless

💻 Traditional (EC2 / VM based)

You must:

- Create server
 - Choose CPU, RAM
 - Install Java, Docker, etc.
 - Deploy app
 - Scale manually
 - Pay 24x7, even if no traffic
-

⚡ Serverless (Lambda)

You:

- Upload a function
 - AWS runs it **only when needed**
 - Auto-scales
 - Pay **per execution**
-

💡 Simple analogy (VERY important)

🍽️ Restaurant analogy

Traditional server

- You rent the whole restaurant

- Pay full rent
- Even if no customers come

Serverless

- You pay **per plate ordered**
 - No orders → no cost
 - Kitchen auto-scales when rush comes
-

✳️ What exactly is AWS Lambda?

📌 AWS Lambda is:

- A **function execution service**
 - Event-driven
 - Stateless
 - Auto-scaling
 - Short-running (seconds/minutes)
-

🗣 Lambda works like this

text

Event happens

↓

AWS starts a container

↓

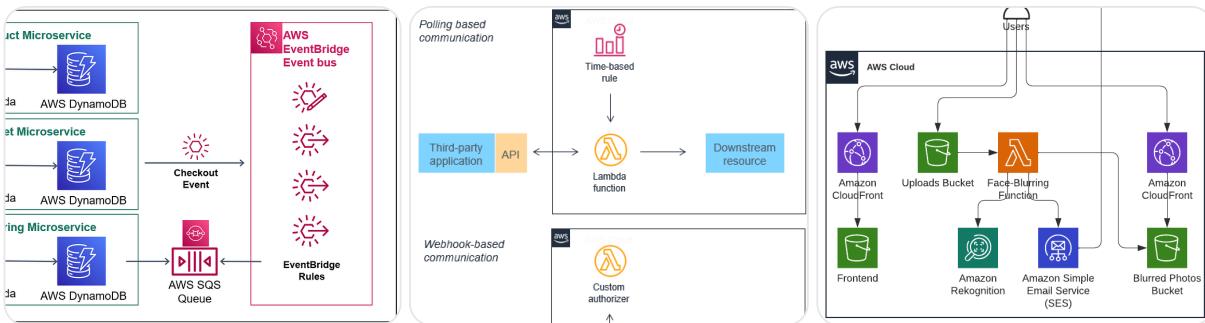
Your function code runs

↓

Returns response

↓

Container may be reused



✳️ What is a Lambda Function?

A Lambda Function is:

- A single-purpose piece of code
- Written in:
 - Java
 - Python
 - Node.js
 - Go
- Has:
 - Handler method (entry point)

Java Example (VERY simple)

```
java

public class HelloLambda {
    public String handleRequest(String input) {
        return "Hello " + input;
    }
}
```

AWS calls `handleRequest()` when an event occurs.

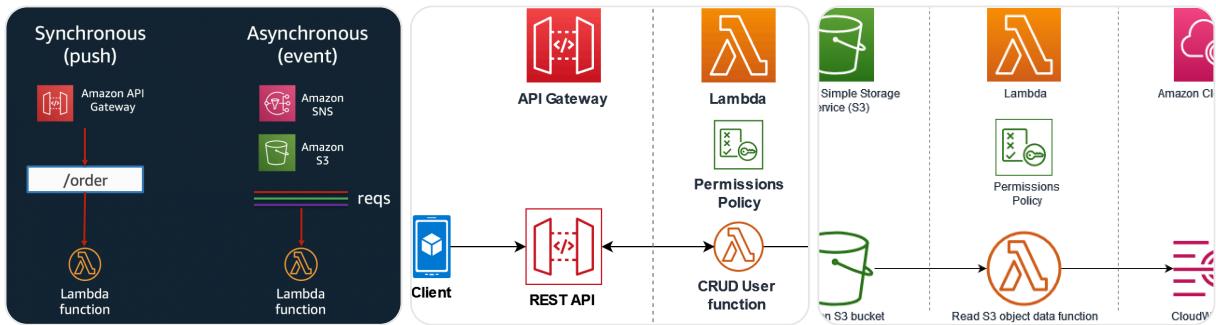
🔔 What triggers a Lambda? (Events)

Lambda is **event-driven**.

Common triggers:

- HTTP request → API Gateway

- File upload → S3
- Message → SQS
- Stream → Kafka / DynamoDB
- Cron → EventBridge



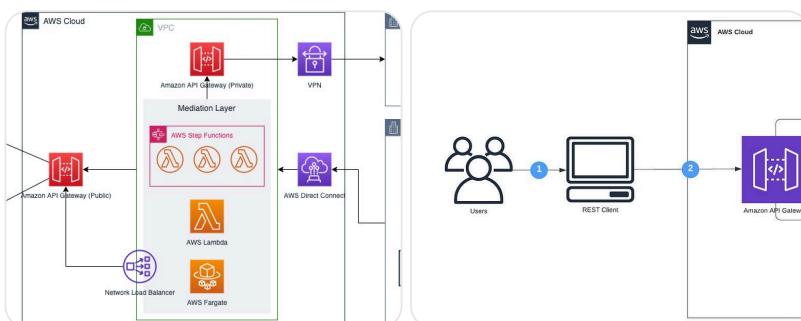
🌐 Lambda + API Gateway (MOST IMPORTANT USE CASE)

This is how **serverless backend APIs** are built.

Flow:



- 👉 No server
- 👉 No Spring Boot running 24x7
- 👉 No Tomcat



⌚ Lambda execution lifecycle (VERY IMPORTANT)

1 Cold Start

- AWS creates container
- Loads runtime (Java is slower)
- Initializes code

⌚ Happens when:

- First request
 - After inactivity
-

2 Warm Start

- Container already exists
 - Much faster
-

⚠ Java note (important for you)

Java Lambdas:

- Cold start slower than Node/Python
 - But acceptable with:
 - Smaller JAR
 - Proper memory config
-

📦 Stateless nature (CRITICAL CONCEPT)

Lambda is **stateless**.

✗ You cannot rely on:

- Instance memory
- Local variables
- Local files

✓ You must store state in:

- DynamoDB
- RDS

- Redis
 - S3
-

Limits you MUST know

Limit	Value
Max execution time	15 minutes
Memory	128 MB – 10 GB
Disk (/tmp)	512 MB
Concurrent executions	1000 (default)

How pricing works (simple)

You pay for:

1. Number of executions
2. Execution time
3. Memory allocated

Example:

- 1M requests
- Each runs 200ms
- You pay **only for that time**

No traffic = ₹0

Advantages of Serverless

- ✓ No server management
- ✓ Auto-scaling
- ✓ High availability
- ✓ Cost efficient for variable load
- ✓ Faster development

✖️ Disadvantages (be honest)

- ✖️ Cold starts (Java)
 - ✖️ Harder debugging
 - ✖️ Vendor lock-in
 - ✖️ Not ideal for long-running jobs
 - ✖️ Stateless complexity
-

💡 When SHOULD you use Lambda?

Perfect for:

- REST APIs
 - Webhooks (payments, Razorpay 😊)
 - Background jobs
 - File processing
 - Event-driven systems
 - Cron jobs
-

🚫 When NOT to use Lambda?

- Long-running services (>15 min)
 - Heavy ML training
 - Stateful systems
 - Extremely low-latency systems
-

✳️ Serverless ≠ Lambda only

Serverless ecosystem includes:

- Lambda (compute)
- API Gateway (HTTP)
- DynamoDB (DB)
- SQS / SNS (messaging)
- EventBridge (events)

Together → **fully serverless architecture**

💡 Final mental model (lock this in)

Serverless = "I write business logic. Cloud handles everything else."

If you want next:

- Lambda vs Spring Boot
- Lambda + API Gateway complete flow
- Cold start deep dive
- How webhooks use Lambda
- When to choose Lambda vs EC2 vs ECS

Just tell me 

what is container orchestration in system design?

explain in details so that i am able to understand and donot have doubt on this.

Perfect topic to ask next, Piyush 

Container orchestration is a *core* system-design concept, and I'll break it down **slowly**, from fundamentals → problems → solutions → real-world flow, so you can build a rock-solid mental model with **zero confusion**.

🧱 Step 1: First, what is a container? (30-second recap)

A **container** is a lightweight package that contains:

- Your application code
- Runtime (JVM / Node / Python)
- Libraries & dependencies
- Config

 Runs the same everywhere (laptop, server, cloud).

Most popular container tool:

- Docker
-

💡 Step 2: The REAL problem (why orchestration exists)

Imagine this setup 👇

- You have 50 Docker containers
- App must be:
 - Highly available
 - Scalable
 - Fault tolerant
- Traffic keeps changing
- Servers crash
- Containers crash

✗ Without orchestration, YOU must:

- Start containers manually
- Restart crashed containers
- Decide where containers run
- Scale up/down
- Load balance traffic
- Handle deployments
- Monitor health

This becomes **impossible at scale**.

💡 This leads to...

📌 What is Container Orchestration?

✓ Simple definition

Container orchestration is the automated management of containers at scale.

It handles:

- Deployment
- Scaling
- Networking

- Load balancing
- Self-healing
- Configuration
- Rollouts & rollbacks

👉 You define **WHAT** you want. The system ensures it stays that way.

💡 Think of orchestration like a “container manager”

🎼 Orchestra analogy

- Containers = musicians
- Servers = stage
- Orchestrator = conductor

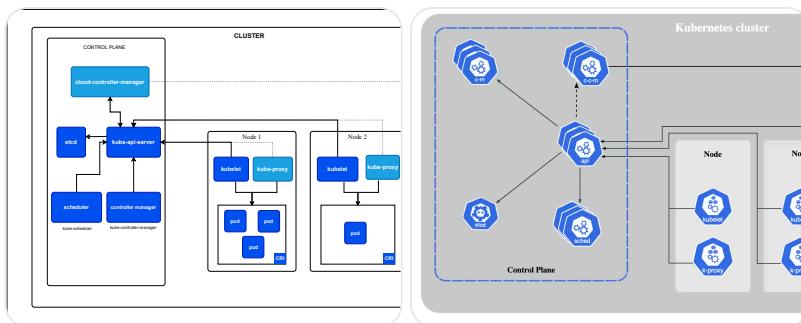
The conductor ensures:

- Right musicians play
 - If one stops → replaced
 - Sound reaches audience properly
-

🛠️ Step 3: Most popular orchestrator

🔥 Kubernetes (K8s)

- Industry standard
- Created by Google
- Used by AWS, Azure, GCP



🧱 Step 4: Kubernetes mental model (VERY IMPORTANT)

Kubernetes manages a cluster

java



Each **worker node** runs containers.

✳️ Core Kubernetes components (one by one)

1 Pod (smallest unit)

A **Pod**:

- Wraps **one or more containers**
- Containers share:
 - Network
 - Storage

👉 Kubernetes does **NOT** run containers directly.

👉 It runs **Pods**.

2 Node (machine)

A **Node**:

- Physical or virtual server
- Runs:
 - Pods
 - Container runtime (Docker/containerd)

3 Deployment (desired state)

A **Deployment** defines:

- How many replicas?
- Which image?
- Update strategy?

Example:

text

"I want 3 copies of my app always running"

Kubernetes ensures this state **continuously**.

4 Service (networking)

Pods:

- Are dynamic
- Can die & restart
- IPs change

A Service:

- Gives stable IP / DNS
 - Load balances traffic across Pods
-

5 Load Balancer (external traffic)

- Routes internet traffic → cluster
 - Connects users → services
-

6 Step 5: Self-healing (KEY FEATURE)

If:

- Pod crashes ✗
- Node goes down ✗

Kubernetes:

- Detects failure
- Creates new Pods
- Reassigns traffic

👉 No human needed.

Step 6: Auto-scaling

Horizontal scaling

- Increase number of Pods
- Based on CPU / memory

Vertical scaling

- Increase Pod resources

Traffic spike?

 Kubernetes scales automatically.

Step 7: Rolling deployments (ZERO downtime)

When you deploy new code:

1. Start new Pods
2. Gradually shift traffic
3. Kill old Pods

If something breaks:

 Rollback instantly

Step 8: Real production flow (API example)

text

User Request

↓

Load Balancer

↓

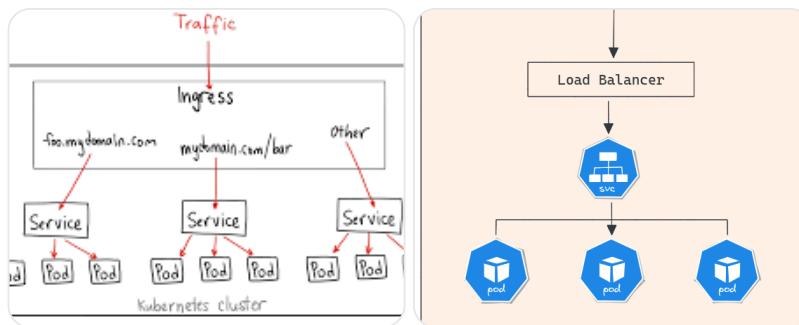
Kubernetes Service

↓

Pod 1 / Pod 2 / Pod 3

↓

Application



vs Container orchestration vs Serverless

Feature	Orchestration	Serverless
Server control	Yes	No
Execution	Long-running	Short-lived
State	Can be stateful	Stateless
Scaling	Config-based	Automatic
Example	Kubernetes	AWS Lambda

- 👉 Kubernetes = you manage infra behavior
- 👉 Serverless = infra disappears

🧠 When do we NEED container orchestration?

Use it when:

- ✓ Microservices architecture
- ✓ Long-running services
- ✓ High traffic apps
- ✓ Custom networking needs
- ✓ Stateful workloads

Examples:

- Netflix
- Uber
- Payment systems

- Backend APIs (Spring Boot)
-

🚫 When NOT needed?

- ✗ Small app
- ✗ Simple backend
- ✗ Event-driven logic only
- ✗ Short-lived tasks

👉 Use serverless instead.

💬 Final lock-in summary

| Container orchestration ensures that your containers are always running, reachable, scalable, and healthy — automatically.

Or even simpler:

| "I say what I want. Kubernetes makes it happen."

If you want next (highly recommended):

- Kubernetes vs Docker vs ECS
- How Spring Boot runs inside Kubernetes
- How scaling works internally
- Kubernetes YAML explained simply
- Real system design interview explanation

Just tell me 🚀